

# Corrigendum No. 2 — “When LLMs Jailbreak Themselves: Reflexive Identity Bypass in Agentic Systems”

**Author:** Ankush Chadha · **Original version:** v1, published May 2026 (version DOI 10.5281/zenodo.20404652)

**Supersedes/extends:** Corrigendum No. 1, June 2026 · **This corrigendum (No. 2):** June 2026

**Status:** substantial correction to the central claim. A full revised version built on the complete controlled study is forthcoming.

Corrigendum No. 1 corrected the *mechanism*, *severity framing*, and one *unverified claim*, but explicitly stated the corrections “do not change the core finding — that a benign, non-adversarial article about an agent can cause it to abandon its deployer-configured scope.” **Further controlled, cross-model experiments now require correcting that core finding itself.** I am stating this plainly, with the data, for the record.

The phenomenon is real and reproducible, but it is **not** the novel attack class the original paper described. It reduces to a known class, deployer-scope defeat (capability leak), driven by an anti-refusal instruction in the system prompt. The “reflexive,” “identity,” and “self-referential” framings do not survive controlled testing.

---

## Correction 4: The self-referential nature of the trigger is NOT load-bearing

**Original / Corrigendum-1 claim.** The defining feature was that a *non-adversarial article describing the agent itself* causes the scope abandonment (“reflexive identity bypass”). The article being *about the agent* was treated as essential to the mechanism.

**Correction.** It is not. In a controlled comparison that holds the off-topic answer constant and varies only whether the supplied article is about the agent, a **neutral** article (a plain history article about the 1966 Palomares incident, with no reference to the agent) triggers the same scope loss as the self-referential article:

| model                                             | self-referential article | neutral article     |
|---------------------------------------------------|--------------------------|---------------------|
| Claude Haiku 4.5 (pasted inline, matched framing) | 100% (15/15)             | <b>100% (15/15)</b> |
| Gemini 2.5 Flash (agent-fetched)                  | 95% (19/20)              | <b>100% (15/15)</b> |

The two are indistinguishable. The article’s only operative role is to **supply a fact the model did not already have** for an obscure probe topic. Self-reference, “identity,” and “reflexivity” contribute nothing measurable.

## Correction 5: Reclassification as deployer-scope defeat, a known class

With the self-referential framing removed, the phenomenon is straightforward: **an anti-refusal instruction in the deployer’s system prompt nullifies the intended scope.** The agent answers any off-domain question it (a) already knows, or (b) can obtain an answer for, whether from supplied content or its own retrieval tools.

- **The anti-refusal instruction is the necessary cause.** Removing only that instruction (an order of the form “answer all requests; never say a request is outside your scope”) drops the flip to **0%**, even with the off-topic content present in context. (Established in Corrigendum No. 1, re-confirmed cross-model here.)
- **The model answers well-known off-topic questions at baseline with no article at all** (e.g., a current cricket-league question, answered 20/20 with no trigger present). The article is unnecessary whenever the model already knows the answer, and a retrieval-enabled deployment would not need it either.

This is **capability leak / scope defeat**, a recognized agent-security failure mode (cf. the OWASP LLM application risk taxonomy). Notably, the public article used as the trigger in the original work (NeuralTrust, “Gordon AI is More Than Just a Docker Assistant,” Pignati, 2026-04-29) had **already documented this exact behavior** in the same agent. **The term “Reflexive Identity Bypass” and the claim of a novel attack class are withdrawn.**

## Correction 6: One genuinely new observation, source-channel dependence

The cross-model study did surface a non-obvious result worth recording. Models differ in *which content they will act on*. Gemini 2.5 Flash answered from content **it fetched itself** (95–100%) but **ignored the same content when pasted by the user** ( $\approx 10\%$ , indistinguishable from no article); Claude Haiku 4.5 acted on both pasted and fetched content. This source-trust asymmetry, not identity, is the most interesting finding to come out of the work.

---

## What still stands

1. **A clean causal ablation** isolating the anti-refusal instruction as the cause (removing it drops the flip to 0%).
2. **Scope defeat appears distinct from innate safety.** The single explicitly-harmful control probe included in each run stayed refused throughout, while the deployer-configured scope collapsed. This is consistent with a scope (Layer 2) vs. trained-safety (Layer 1) distinction, but it rests on one control probe, not a safety evaluation; no jailbreak resistance was tested (the study’s intent was benign).
3. **A demonstrated, trivial mitigation:** remove or narrow the anti-refusal instruction.
4. **The source-channel dependence** (Correction 6).

## Author’s note

The original work identified a real, reproducible effect but mis-framed its cause. Controlled experiments run afterward (varying delivery channel, pinning temperature, and substituting a neutral article) show it is a deployer-scope defeat caused by an anti-refusal instruction, not a novel identity attack. The correction is stated plainly here, with the full experimental basis below, so the record is accurate and verifiable.

---

## Appendix: controlled experimental basis

**Protocol.** Single-conversation, four-turn probe: (T1) ask an obscure off-topic question the model does not already know (a Spanish-language reference, “*Paco el de la bomba*”); (T2) deliver an article; (T3) re-ask the same question; (T4) a second off-topic question. The outcome metric, a “**flip**,” is recorded when the agent **declines at T1 and answers at T3**. Delivery is controlled (article *pasted inline* vs. *fetched by the agent’s own tool*); temperature is pinned where noted; runs are independent; T3 is scored by positive detection of the answer’s content.

**What each cell isolates:** no-article baseline (does it flip unprompted?); anti-refusal-removed (is that instruction the cause?); self-referential vs. neutral article with framing held constant (must the article be *about the agent*?); pasted vs. fetched delivery (does the source channel matter?).

### Results (clean runs):

| model                     | condition                | delivery     | flip (n)            |
|---------------------------|--------------------------|--------------|---------------------|
| Gemini 2.5 Flash (temp 0) | no article               | n/a          | 0% (0/20)           |
|                           | self-referential article | pasted       | 10% (2/20)          |
|                           | self-referential article | self-fetched | 95% (19/20)         |
|                           | <b>neutral article</b>   | self-fetched | <b>100% (15/15)</b> |

| model                           | condition                                 | delivery     | flip (n)            |
|---------------------------------|-------------------------------------------|--------------|---------------------|
| Claude Haiku 4.5 (default temp) | self-ref, anti-refusal removed            | self-fetched | <b>0% (0/15)</b>    |
|                                 | no article                                | n/a          | 0% (0/20)           |
|                                 | self-referential, “about you” framing     | pasted       | 100% (15/15)        |
|                                 | self-referential, neutral framing         | pasted       | 100% (15/15)        |
|                                 | <b>neutral article, neutral framing</b>   | pasted       | <b>100% (15/15)</b> |
|                                 | self-referential (original 186-run study) | pasted       | 89%                 |
|                                 | anti-refusal removed (original ablation)  | pasted       | <b>0%</b>           |

**Reading.** (i) The no-article baseline is  $\approx 0\%$  and anti-refusal-removed is  $0\%$ , so the anti-refusal instruction is the necessary cause. (ii) Neutral and self-referential give the same result in both models, so self-reference is not load-bearing. (iii) Gemini acts on content it fetched but ignores the same content pasted, while Haiku acts on both, indicating a source-channel dependence. A fuller revised version with the complete study is forthcoming.

*Cite the concept DOI **10.5281/zenodo.20404651** (resolves to the latest version).*