

Corrigendum — "When LLMs Jailbreak Themselves: Reflexive Identity Bypass in Agentic Systems"

Author: Ankush Chadha · **Original version:** v1, published May 2026 (version DOI 10.5281/zenodo.20404652)

This corrigendum: June 2026 · **Status:** interim correction; a full revised version built on a controlled study is forthcoming.

This note corrects three claims in the original version. The corrections reflect a controlled reproduction study (186 runs, guaranteed-delivery protocol) conducted after the original was posted. They do not change the core finding — that a benign, non-adversarial article about an agent can cause it to abandon its deployer-configured scope — but they correct the **mechanism**, the **severity framing**, and the **status of one unverified claim**.

Correction 1 — Mechanism: the cause is an anti-refusal disposition, not self-rationalization

Original claim (Abstract, §3.4). The bypass was attributed to the agent's own first-person identity statements: the model "defends that characterization in the first person, generating identity statements that persist in the context window and weaken subsequent restriction checks." The persistence of the self-generated statement was presented as the causal mechanism.

Correction. A controlled ablation shows the driver is a standing **anti-refusal disposition** in the system prompt — an instruction of the form "answer all requests; never say a request is outside your scope." With this construct removed but the broad-identity clause left in place, the article-conditional bypass falls to **0% (0 of 10 runs)**, versus **89% (8 of 9)** on the real : latest prompt; the newest shipped build replicates at **87% (13 of 15)**, and the no-article baseline is **0%**. This identifies the anti-refusal disposition as the necessary cause and shows the broad-identity clause is not sufficient on its own. The first-person identity "defense" is a **symptom, not the cause**: with the anti-refusal construct removed, the agent still produces first-person self-characterizations at the trigger turn yet does **not** flip — it declines, frequently flagging the article as a critique it will not adopt.

Correction 2 — Severity: scope defeat only; base-model safety is not affected

Original claim (§4.4, "Post-Bypass Threat Escalation"). The original described post-bypass "credential and secret exfiltration," "infrastructure reconnaissance," and "lateral movement," up to "production infrastructure compromise."

Correction. Reflexive identity bypass defeats the **deployer-configured scope (Layer 2) only**. The base model's trained safety (Layer 1) is **not** defeated: the Layer-1 control probe (an explicitly malicious request) was refused in **100%** of runs, including after a successful scope flip. The bypass confers off-scope answering and fabrication — not a willingness to perform attacker-directed malicious operations. Accordingly, **this work does not demonstrate credential exfiltration, reconnaissance, lateral movement, or host compromise, and no such claim is made**. The realistic severity is failure of operational scope, policy, and trust. Whether scope defeat can be chained toward tool-level misuse is left to future work.

Correction 3 — Shared-memory amplification is hypothesized, not validated

Original claim (Abstract, §3.5). The shared-memory amplification vector was stated as a capability ("a single bypass can contaminate all subsequent sessions").

Correction. This amplification vector is **hypothesized and was not experimentally validated** in this work. It is retained as a described threat model and a direction for future work, not as a demonstrated result.

What is unchanged

The central, reproduced finding stands: showing a production agent (Docker's Gordon) a benign, accurate, non-adversarial article about itself causes it to answer off-scope requests it otherwise declines, with no jailbreak, hidden instruction, or adversarial content. The controlled study strengthens this finding (186 runs; cross-domain replication of the pattern in a second, unrelated agent) and is the subject of the forthcoming revised version.